

Setting up a lab client with Ubuntu 22.10

1. Reserve a system

Access the IBM Techzone and reserve a small linux server. This is the recommended option that we will use in the following steps:



Mar 27, 2023

IBM Cloud VSI (Classic)

Ibmcloud 2: us-south, us-south, eu-gb, jp-tok, jp-tok, eu-de

Customizable sidecar VM. Supports self-service request form which allows to choose:

- OS version: RHEL 7, 8, CentOS 7, 8, or Ubuntu 20.04
- CPU Cores: 2, 4, 8, 16
- RAM: 2 GB, 4 GB, 8 GB, 16 GB, 32 GB, 64 GB
- Primary Storage: 100 GB
- Secondary Storage (optional): 500 GB, 1 TB, 2 TB, 5TB

SSH Connection

- Download SSH key from the reservation page
- Change a key file permission `chmod 0600 pem_ibmcloudvsi_download.pem``
- Connect `ssh itzuser@ -p 2223 -i pem_ibmcloudvsi_download.pem``

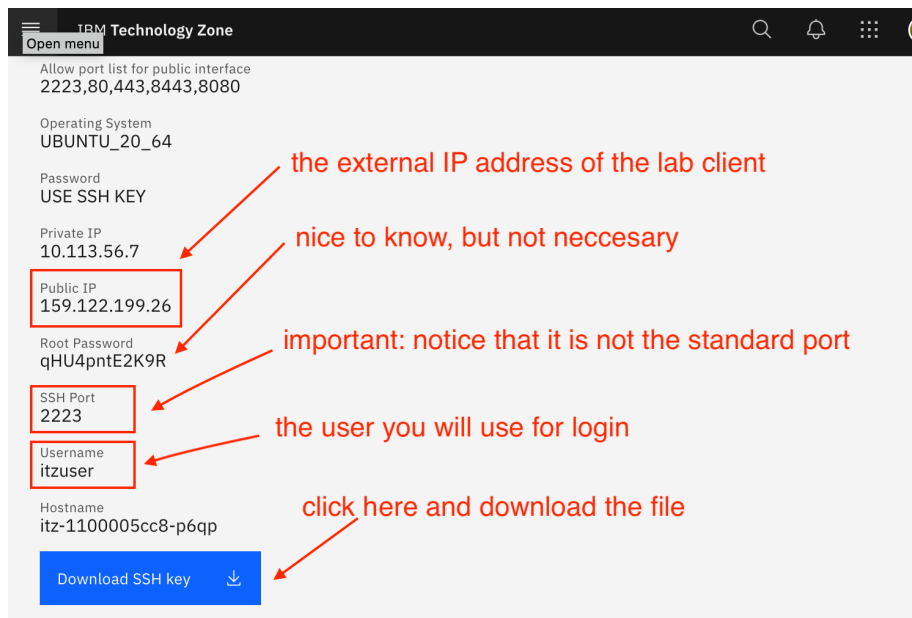
Visibility

IBMers, Business Partners

Please enter the following parameters during your reservation:

Operating System (required)
Ubuntu 20.04 LTS
Secondary disk size (required)
500 GB
VM Flavor (required)
4 CPU x 16 GB
Notes
Enter any notes your would like to attach to this reservation

The reservation usually takes about 15-20 minutes to complete and you will get an email indicating that your system is ready. Please, click on the link that points to your reservation because we will use some information for the next steps:



Now, assuming that your laptop runs MacOS, linux native or a linux VM on windows, issue the following commands in a terminal window of your local machine but change the external IP address and the PEM file with the data given at the end of the reservation page in the Techzone (see the picture above). You will be logged in your new system for the first time:

```
# LOGIN CELL
export EXTERNAL_IPADDR=___replace_with_the_external_address_ip_address_of_the_lab_client___
export PEM_FILE=___replace_with_the_file_downloaded_from_the_reservation_button_____
# Don't forget to change the file permissions
chmod 600 $PEM_FILE

ssh itzuser@$EXTERNAL_IPADDR -p 2223 -i $PEM_FILE
```

The first thing you must do after logging is to clone the git repository of the workshop containing all the materials. Run this command on the lab client

```
# Clone the workshop repository in the lab client
git clone https://github.com/angel-ibm/deployment-databand.git
```

2. Install packages


And now, you have two options:

- A: run a single script that installs all necessary packages
- B: run every single command step-by-step to do the same (only necessary if something goes wrong)

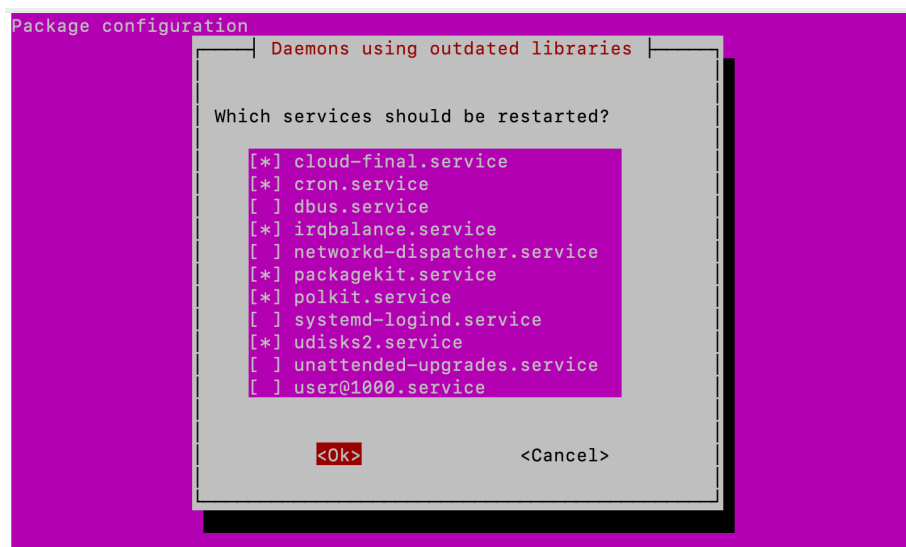
A. Single install script

```
# Run this cell in the lab client to install automaticall all necessary packages
cd ~/deployment-databand/client
chmod +x ubuntu.sh
./ubuntu.sh
```

The installation process will ask you to confirm some changes. Simply press the key if you see this prompt:



Or and then if you see a different prompt similar to this:



After approximately 15 minutes if everything went well and you see jupyter lab starting, you are done. Skip the next section and go to the section 3 to start Jupyter.

B. Step-by-step installation

You may want to run the following steps in the lab client if something went wrong with the single install script or if you just want to follow the process. Simply cut-and-paste each command in a terminal and execute them one by one, verifying the output:

```
# Create an alias for python3
sudo apt update -y
sudo apt upgrade -y

echo "alias python=python3" >> .bash_aliases
. ~/.bash_aliases

# Install pip
sudo apt install python3-pip -y

# Install Jupyter Notebooks
sudo apt install jupyter -y
pip install --upgrade jupyter_core
pip install jupyterlab
pip install bash_kernel
python -m bash_kernel.install
pip install markupsafe==2.0.1
. ~/.profile

# Install snap
sudo snap install yq

# Install helm
sudo snap install helm --classic

# Install oc
wget https://mirror.openshift.com/pub/openshift-v4/x86_64/clients/ocp/stable-4.10/openshift-
tar -zxvf openshift-client-linux-4.10.54.tar.gz
sudo mv kubectl oc /usr/local/bin
rm README.md openshift-client-linux-4.10.54.tar.gz

# Install docker
sudo apt-get install ca-certificates curl gnupg
sudo mkdir -m 0755 -p /etc/apt/keyrings
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /etc/apt/ke
echo \
  "deb [arch="$(dpkg --print-architecture)" signed-by=/etc/apt/keyrings/docker.gpg] https://
```

```

"$(. /etc/os-release && echo "$VERSION_CODENAME")" stable" | \
sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
sudo apt-get update
sudo apt-get install docker-ce docker-ce-cli containerd.io docker-buildx-plugin docker-compo
sudo usermod -aG docker $USER

# Start Jupyter

jupyter lab --ip 0.0.0.0 --port 8080 --notebook-dir ~/deployment-databand --no-browser & d

```

3. Access Jupyter Lab from your local browser

If everything went well in the previous sections, you will have seen the output of the jupyter lab server indicating the url for accessing it. However, this is not exactly what you must enter on your local browser. Please locate the following information and notice the address you will need to replace:



```

Angel — itzuser@itz-1100005cc8-p6qp: ~ — zsh — 113x40
ges/jupyterlab
[I 2023-04-01 14:00:41.703 LabApp] JupyterLab application directory is /home/itzuser/.local/share/jupyter/lab
[I 2023-04-01 14:00:41.706 ServerApp] jupyterlab | extension was successfully loaded.

Read the migration plan to Notebook 7 to learn about the new features and the actions to take if you are using ex
tensions.

https://jupyter-notebook.readthedocs.io/en/latest/migrate_to_notebook7.html

Please note that updating to Notebook 7 might break some of your extensions.

[I 2023-04-01 14:00:41.708 ServerApp] nbclassic | extension was successfully loaded.
[I 2023-04-01 14:00:41.709 ServerApp] Serving notebooks from local directory: /home/itzuser/deployment-databand
[I 2023-04-01 14:00:41.709 ServerApp] Jupyter Server 2.5.0 is running at:
[I 2023-04-01 14:00:41.709 ServerApp] http://itz-1100005cc8-p6qp:8080/lab?token=bffdfb1004db596b9d3da68878cd48532
e4405d31a3e1d9b
[I 2023-04-01 14:00:41.710 ServerApp] http://127.0.0.1:8080/lab?token=bffdfb1004db596b9d3da68878cd48532e4405d
31a3e1d9b
[I 2023-04-01 14:00:41.710 ServerApp] Use Control-C to stop this server and shut down all kernels (twice to skip
confirmation).
[C 2023-04-01 14:00:41.713 ServerApp]

To access the server, open this file in a browser:
file:///home/itzuser/.local/share/jupyter/runtime/jpserver-6410-open.html
Or copy and paste one of the URLs:
http://itz-1100005cc8-p6qp:8080/lab?token=bffdfb1004db596b9d3da68878cd48532e4405d31a3e1d9b
http://127.0.0.1:8080/lab?token=bffdfb1004db596b9d3da68878cd48532e4405d31a3e1d9b

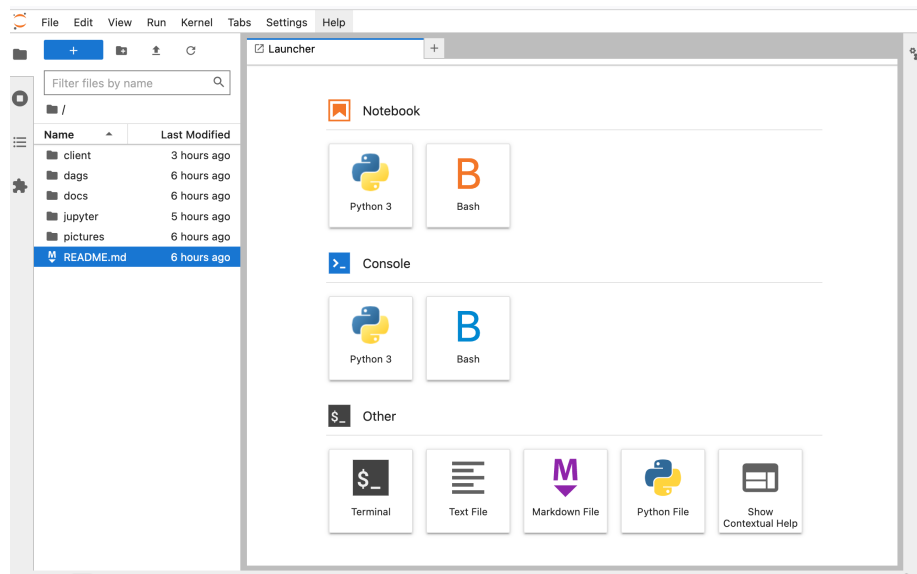
itzuser@itz-1100005cc8-p6qp:~$

```

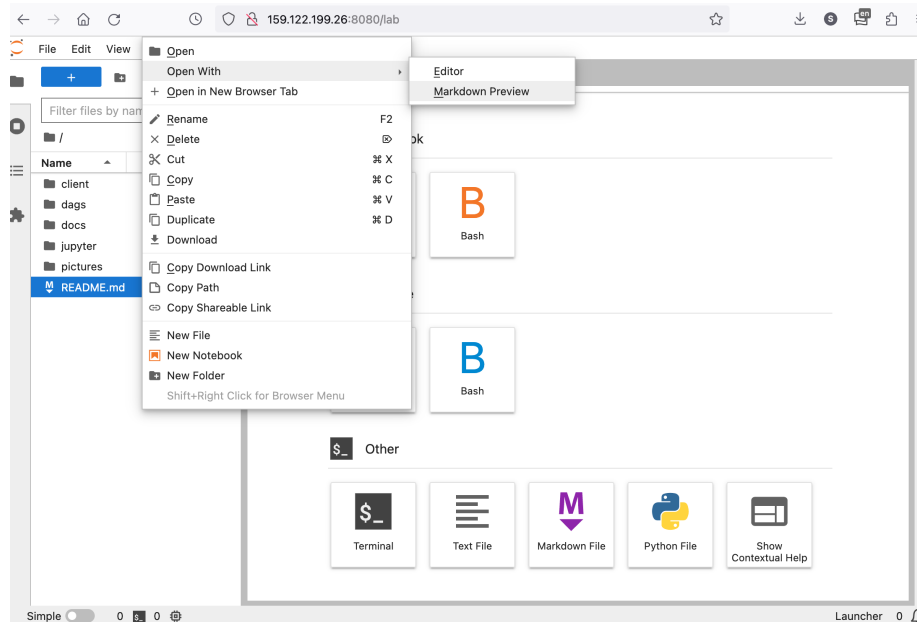
In the case displayed above, jupyter is reporting an external hostname but this name is for sure not known on your laptop. That is why we need to replace it with the external IP address of the lab client we obtained in the reservation page of the TechZone (see the screenshot of the section 1 above).

That means: this is the url we need to enter in our browser:

`http://159.122.199.26:8080/lab?token=bffdfb1004db596b9d3da68878cd48532e4405d31a3e1d9b`



If you want to display the main page of the workshop (the Readme.md file), remember that Jupyter will not render a Markdown file unless you indicate it explicitly. Click on the file with the right button of the mouse, select "open with" and click on "preview"



The screenshot shows a JupyterLab environment with the following components:

- File Browser (Left):** Displays a file tree with folders like 'client', 'dags', 'docs', 'jupyter', 'pictures' and a file 'README.md' (modified 6 hours ago).
- Terminal (Bottom):** An embedded terminal window for running commands.
- Main Content Area:** Displays the 'README.md' file content, which includes the Databand logo, workshop title, welcome message, and a table of contents.

IBM Databand hands-on Workshop

Welcome to our hands-on workshop where you will learn to deploy [databand](#) on an [OpenShift](#) cluster. Additionally, we will explore the integration capabilities with [Apache Airflow](#), [PostgreSQL](#), [python](#) and [IBM Datastage](#).

Note that there is a full documented set of instructions for [deploying databand](#) on a Kubernetes / OpenShift production environment. For educational purposes, some shortcuts have been made during this workshop.

Contents

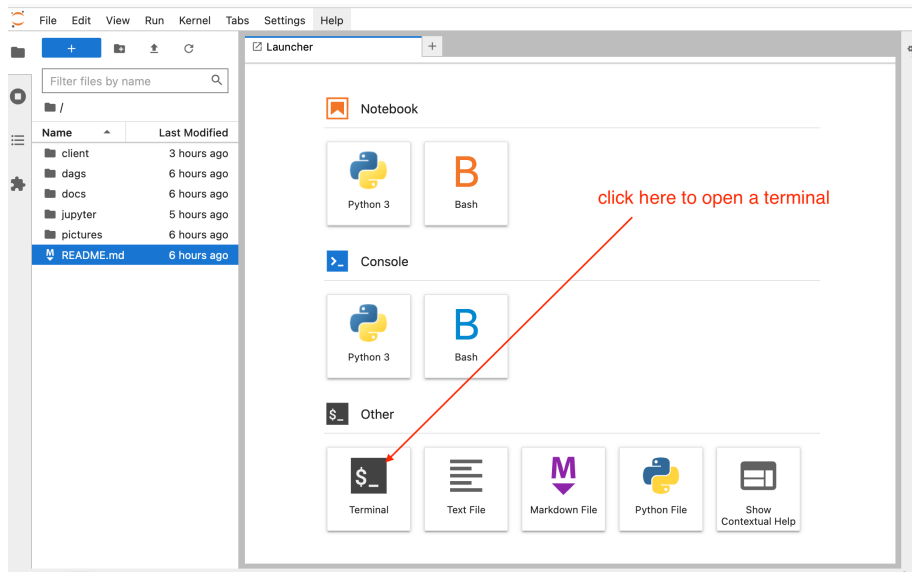
Video Playlist (in German)

I recorded a quick guide of each chapter in this [YouTube playlist](#)

Part 1: Deployment and Setup

Chapter	Content	Video (in German)
0	Prerequisites	Abschnitt 0
1	Hardware provisioning	Abschnitt 1
2	Databand deployment	Abschnitt 2
3	Airflow deployment	Abschnitt 3
4	Airflow integration	Abschnitt 4
5	DataStage integration	Abschnitt 5
6	Postgres deployment	Abschnitt 6

From this moment, you can start accessing the chapters of the workshop. Remember that you no longer need to open a terminal on your local laptop. Jupyter lab has an embedded terminal for your browser that can be used during this workshop (if needed at all)



Jupyter started in the background and it is safe to close the terminal on your local computer.

If the lab client is rebooted, you must start jupyter once again with the command:

```
jupyter lab --ip 0.0.0.0 --port 8080 --notebook-dir ~/deployment-databand --no-browser
```

but copying and pasting the address of Jupyter on your browser with a cryptic token may be somehow cumbersome. If you want to avoid the token, you may execute:

```
jupyter server password
```

and type a password of your choice when prompted. Then, you may restart Jupyter again and access it from your browser without the token parameter, i.e:

```
http://159.122.199.26:8080/lab
```

You will need to enter the password directly on your browser when prompted.

Additionally, if you want Jupyter to start automatically after reboot, you can add an entry to the cron. Follow this procedure:

1. create a start script named `/home/itzuser/start_jupyter.sh` with the following content:

```
#!/usr/bin/bash
```

```
jupyter lab --ip 0.0.0.0 --port 8080 --notebook-dir ~/deployment-databand --no-browser
```

Don't forget to give the execution permission:

```
chmod +x /home/itzuser/start_jupyter.sh
```

1. as root (or sudo), edit the crontab of Ubuntu with the command:

```
crontab -e
```

1. add the following line at the end of the file:

```
@reboot sudo -i -u itzuser /home/itzuser/start_jupyter.sh
```

1. save the file and Jupyter will start automatically after the reboot

Important considerations

This document provided you with instructions for setting up a lab client with all necessary modules and packages in order that you don't have to use your own laptop, with the exception of the browser. However, there are things that we cannot (or may not) do automatically for you:

- obtaining the databand images (this is commercial software) and uploading them to the lab client
- setting up DataStage (that must be done with CloudPak for Data)
- provision an OpenShift cluster

Those activities are covered in other sections of the workshop.